



Previous note

The *sofs18* is a simple and limited file system, based on the ext2 file system, which was designed for purely educational purposes and is intended to be developed in the practical classes of the Operating Systems course in academic year of 2018/2019. The physical support is a regular file from any other file system.

1 Introduction

Almost all programs, during their execution, produce, access and/or change information that is stored in external storage devices, which are generally called mass storage. Fit in this category magnetic disks, optical disks, SSD, among others.

Independently of the physical support, structurally one can verify that:

- mass storage devices are usually seen as an array of blocks, each block being 256 to 8 Kbytes long;
- blocks are sequentially numbered (LBA model), and access to a block, for reading or writing, is done given its identification number.

Direct access to the contents of the device should not be allowed to the application programmer. The complexity of the internal structure and the necessity to enforce quality criteria, related to efficiency, integrity and sharing of access, demands the existence of a uniform interaction model. The concept of **file** appears, then, as the logic unit of mass memory storage. This means reading and writing in a mass storage device is always done in the context of files.

From the application programmer point of view, a file is an abstract data type, composed of a set of attributes and operations. Is the operating system's responsibility to provide a set of system calls that implement such abstract data type. These system calls should be a simple and safe interface with the mass storage device. The component of the operating system dedicated to this task is the **file system**. Different approaches conduct to different types of file systems, such as NTFS, ext3, FAT*, UDF, APFS, among others.

1.1 File as an abstract data type

The actual attributes of a file depend on the implementation. The following represents a set of the most common ones:

name — an identifier used to access the file upon creation;

internal identifier — a unique (numerical) internal identifier, more suitable to access the file from the file system point of view;

size — the size in bytes of the file's data;

ownership — an indication of who the file belongs to, suitable for access control;

permissions — an attribute that in conjunction with the ownership allows to grant or deny access to the file;

monitoring times — typically, time of last access and time last modification;

localization of the data — means to identify the blocks where the file's data is stored;

type — the type of files that can be hold by the file system; here, 3 types are considered:

regular file — what a user usually define as a file;

directory — an internal file type, with a predefined format, that allows to view the file system as a hierarchical structure of folders and files;

shortcut (symbolic link) — an internal file type, with a predefined format, that points to the absolute or relative path of another file.

The operations that can be applied to a file depend on the operating system. However, there is a set of basic ones that are always present. These operations are available through system calls, i.e., functions that are the entry points into the operating system. It follows a list, not complete, of the system calls provided by Linux to manipulate the 3 considered file types:

- common to the 3 file types: `open`, `close`, `chmod`, `chown`, `utime`, `stat`, `rename`;
- common to regular files and shortcuts: `link`, `unlink`;
- only for regular files: `mknod`, `read`, `write`, `truncate`, `lseek`;
- only for directories: `mkdir`, `rmdir`, `getdents`;
- only for shortcuts: `symlink`, `readlink`;

You can get a description of any one of these system calls executing, in a terminal, the command

```
man 2 <syscall>
```

where `<syscall>` is one of the aforementioned system calls.

1.2 FUSE

In general terms, the introduction of a new file system in an operating system requires the accomplishment of two different tasks. One, is the integration of the software that implements the new file system into the operating system's kernel; the other, is its instantiation on one or more disk devices using the new file system format.

In monolithic kernels, the integration task involves the recompilation of the kernel, including the software that implements the new file system. In modular kernels, the new software should be compiled and linked separately and attached to the kernel at run time. Any way, it is a demanding task, that requires a deep knowledge of the hosting system.

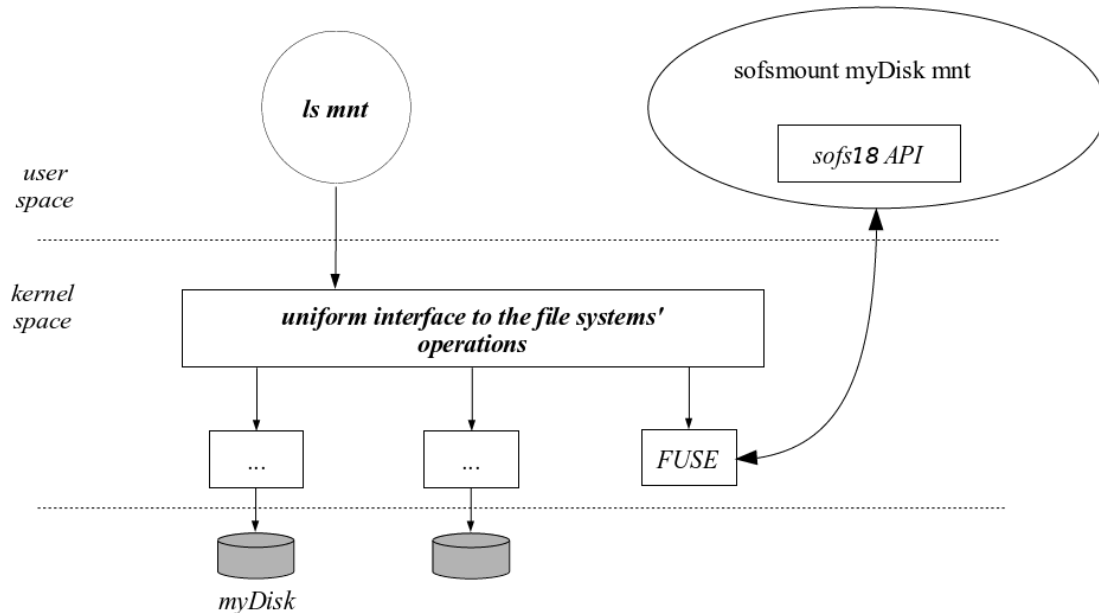
FUSE (File system in User Space) is a canny solution that allows for the implementation of file systems in user space (memory where normal user programs run). Thus, any effect of flaws of the supporting software are restricted to the user space, keeping the kernel immune to them.

The infrastructure provided by FUSE is composed of two parts:

- Interface with the file system — works as a mediator between the kernel system calls and their implementation in user space.


- Implementation library — provides the data structures and the prototypes of the functions that must be developed; also provides means to instantiate and integrate the new file system.

The following diagram illustrates how the *sofs18* file system is integrated into the operating system using FUSE.



2 The *sofs18* architecture

As mentioned above a disk is seen as a set of numbered blocks. For *sofs18* it was decided that each block is 512 bytes long. In a general view, the N blocks of a *sofs18* disk are divided into 5 areas, as shown in the following figure.



Superblock

The **superblock** is a data structure stored in block number 0, which contains global attributes for the disk as a whole and for the other disk structures.

Inodes

An **inode** is a data structure that contains all the attributes of a file, except the name. There is a contiguous region of the disk, called **inode table**, reserved for storing all inodes. This means that the identification of an inode can be given by the index representing its relative position in the inode table. This also means that the number of inodes in a *sofs18* disk is fixed after formatting.

At any given time, there are inodes in use, while others are available to be used (free). When a new file is to be created, a free inode must be assigned to it. It is therefore necessary to:

- define a policy to decide which free inode should be used when one is required;
- define and store in the disk a data structure suitable to implement such policy.

In *sofs18* a FIFO policy is used, meaning that the first free inode to be used is the oldest one in the list. The implementation is based in a list of references to free inodes, stored in a specific region of the disk and described below.

Data blocks

The actual data of any file is stored in blocks taken from the **data block zone**, which is a contiguous sequence of blocks at the end of the disk. This means that the identification of data blocks can be given by the index representing their relative position in the data block zone. This also means that the number of blocks in a *sofs18* disk is fixed after formatting.


Similar to what was stated for inodes, at any given moment, part of the blocks will be in use (by the files stored in the disk), while the others will be available (free). Thus, again, it will be necessary:

- to define a policy to decide which free block should be used when one is required;
- to define and store in the disk a data structure suitable to represent the blocks' states and implement such policy;

In *sofs18*, like for inodes, a FIFO policy is used, meaning that the first free data block to be used is the oldest one in the list. The implementation is based in a list of references to free data blocks, stored in a specific region of the disk and described below.

List of free inodes / data blocks

The list of free inodes and the list of free data blocks have similar implementations. Each list is implemented using a set of contiguous blocks on the disk plus two caches in the superblock, the caches being used to improve the performance of file system accesses. This implementation is illustrated in the next figure.



The table (set of contiguous blocks) represents an array of cells, each one holding the reference to a free inode or a free data block. The cells in the table are managed in a global way. This means that, if a table comprises N blocks and there are R_{PB} references per block, the indices of the cells range from 0 to $N * R_{PB} - 1$. The table is also managed in a circular way, thus, the position after the last one is the first one. Two fields in the superblock indicate the first

occupied cell (head) and the first empty cell (tail). In the figure, the part of the table containing valid references is at gray, while the part at white represents empty cells (containing a special value denoted as null reference).

The insertion cache is used to store references of recently freed items (inodes or data blocks). Thus they represent the tail references in the list. When the insertion cache is full, it must be depleted to the corresponding table. The retrieval cache is used to store the head references in the list. If a retrieval cache is empty, it must be replenished from the corresponding table. The `idx` attribute of the retrieval cache points to the first valid reference in the list. The next elements are those in the cache until its end. After them, follow the elements in the table from `head` to `tail-1`. Finally, it follows the elements in the insertion cache from position 0 to position `idx-1`. Attribute `idx` of the retrieval cache points to the first empty cell.

2.1 List of blocks used by a file (inode)

Blocks are not shared among files, thus, an in-use block belongs to a single file. The number of blocks required by a file to store its information is given by

$$N_b = \text{roundup}\left(\frac{\text{size}}{\text{BlockSize}}\right)$$

where `size` and `BlockSize` represent, respectively, the size in bytes of the file and the size in bytes of a block.

N_b can be very high. Assuming, for instance, that the block size in bytes is 512, a 2 GByte file needs 4 million blocks to store its data. But, N_b can also be very small. In fact, for a 0 bytes file, N_b is equal to zero. Thus, it is impractical that all the blocks used by a file are contiguous in disk. The data structure used to represent the sequence of blocks used by a file must be flexible, growing as necessary.

The access to a file's data is in general not sequential, but instead random. Consider for instance that, in a given moment, one needs to access byte index j of a given file. What is the block that holds such byte? Dividing j by the block size in bytes, one get the index of the block, in the file point of view, that contains the byte. But, what is the number of the block, in the disk point of view? The data structure should allow for a efficient way of finding that number.

In *sofs18*, the defined data structure is dynamic and allows for a quick identification of any data block. Each inode allows to access a dynamic array, denoted d , that represents the sequence of blocks used to store the data of the associated file. Being `BlockSize` the size in bytes of a block, $d[0]$ stores the number of the block that contains the first `BlockSize` bytes, $d[1]$ the next `BlockSize` bytes, and so forth.

Array d is not stored in a single place. The first 4 elements are directly stored in the inode, in the field named `d`. The next elements, when they exist, are stored in an indirect and double indirect way. Inode field `i1` represents the first two elements of an array i_1 ($i_1[0]$ and $i_1[1]$), each element used to indirectly extend array d . Being `RPB` the number of references to blocks that can be stored in a block, $i_1[0]$ is the number of the block that extends array d from $d[4]$ to $d[\text{RPB} + 3]$ and $i_1[1]$ is the number of the block that extends array d from $d[\text{RPB} + 4]$ to $d[2 * \text{RPB} + 3]$.

For bigger files, the double indirect approach is used. Inode field `i2` represents a two-elements array i_2 ($i_2[0]$ and $i_2[1]$), each element used to indirectly extend array i_1 . Thus, $i_2[0]$ is the number of the block that extends array i_1 from $i_1[2]$ to $i_1[\text{RPB} + 1]$ and $i_2[1]$ is the number of the block that extends array i_1 from $i_1[\text{RPB} + 2]$ to $i_1[2 * \text{RPB} + 1]$.

Pattern `NullReference` is used to represent a non-existent block. For example: if $d[1]$ is equal to `NullReference`, the file does not contains block index 1; if $i1[0]$ is equal to `NullReference`, it means $d[4]$ to $d[\text{RPB} + 3]$ are equal to `NullReference`; if $i2[0]$ is equal to `NullReference`, it means $i_1[2]$ to $i_1[\text{RPB} + 1]$ are equal to `NullReference`, and thus $d[2 * \text{RPB} + 4]$ to $d[\text{RPB}^2 + 2 * \text{RPB} + 3]$

are equal to `NullReference`.

2.2 Directories

As stated before, a directory is a set of directory entries. In *sofs18*, a directory entry is a data structure composed of a fixed-size array of bytes, used to store the name of a file, and a reference, used to indicate the inode associated to that file. Thus, a directory can be seen as an array of fixed-size slots, each one with capacity to store a directory entry. The data structure is defined such that a block holds an integer number of directory entries.

The slots of a directory can be in one of two states: in-use or free. A slot is in-use if it contains a directory entry, the name and inode number of an existing file (regular file, directory or shortcut). A slot in the free state has the name field totally filled with the null character (`'\0'`) and the reference field equal to `NullReference`. When a directory entry is deleted, the slot is put in the free state. When a new directory entry is created, the first slot in the free state must be used.

When a directory is created, a data block is assigned to it, with all the slots put in the free state, except the first two. The first two entries have special meaning and are presented in every directory. They are named, `"."` e `".."`, the former representing the directory itself and the latter representing the parent directory. When all the slots of a directory are in-use and a new entry is to be created, a new block is added to a directory, with the entry in the first slot and all the others put in the free state. Thus, the size of a directory is always a multiple of the size of a block. Also, and because of the way deletion is done, a directory never shrinks.

3 Formatting

The formatting operation must fill the blocks of a disk in order to make it an empty *sofs18* device. It is the operation to be performed before the disk can be used.

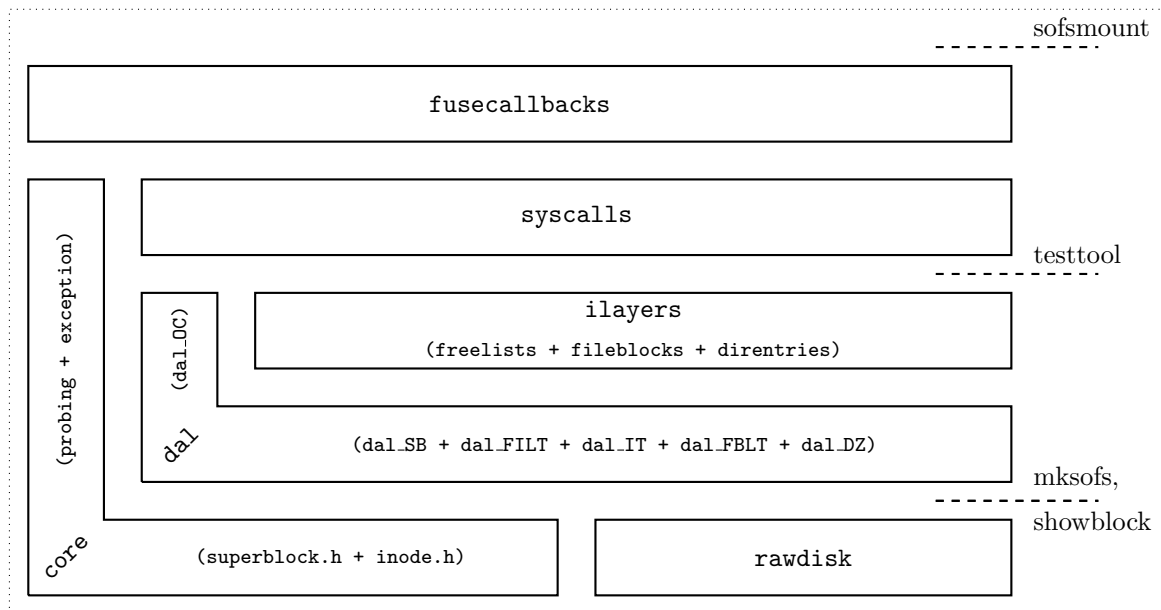
An empty disk (just formatted) has root directory, which in turn has two entries filled, the `"."` e `".."` entries. By convention, the parent of the root is the root itself, thus both entries `"."` and `".."` points to inode number 0.

The forming operation must:

- Choose the appropriated values for the number of inodes, the number of blocks and the number of blocks used by the lists of free inodes and free data blocks, taking into consideration the number of inodes requested by the user and the total number of blocks of the disk. Note that all blocks of the disk must be used, so those that can not be used for other purposes should be added to the inode table.
- Fill in all fields of the superblock, taking into consideration the state of the disk after formatting. The four caches should be put empty.
- Fill in the table of inodes, knowing inode number 0 is used by the root and that all other inodes are free.
- Fill in the list of free inodes table, starting from position 0. The list of free inodes must start in inode number 1 and go sequentially to the last.
- Fill in the list of free data blocks table, starting from position 0. The list of free data blocks must start in data block number 1 and go sequentially to the last.
- Fill in the root directory, knowing that it should occupy block number 0.
- fill in with zeros all free blocks, if such is required by the formatting command.

4 Code structure

The following diagram represents the structure of the *sofs18* code.



rawdisk – This layer implements the physical access to the disk.

core – This layer implements a debugging library (probing), the exception handling (exception) and defines the *sofs18* data types (superblock and inode).

dal – dal stands for disk abstraction layer and implements the access to the different parts of a *sofs18* disk (superblock, free inode list table, inode table, free block disk table and data block zone); it also includes functions to open and close the disk.

ilayers – This layer implements a set of intermediate functions that facilitates the implementation of the system calls. It is composed of 3 modules: management of the free lists (freelists); access to the blocks of a file (fileblocks); and manipulation of directories (direntries).

syscalls – *sofs18* version of file system calls.

fusecallbacks – Interface with FUSE.