



Operating Systems / Sistema de Operação


Interprocess communication

Artur Pereira / António Rui Borges

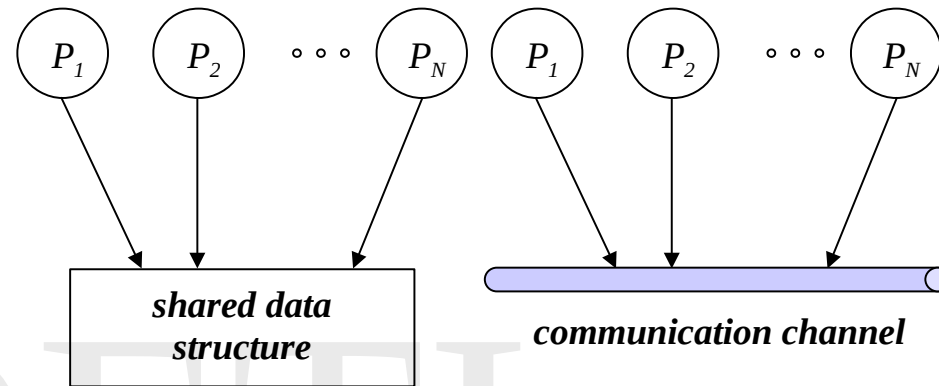
Concepts

- ♦ In a multiprogrammed environment, two or more processes can be:
 - ♦ *independent* – if they, from their creation to their termination, never explicitly interact
 - ♦ actually. there is an implicit interaction, as they compete for system resources
 - ♦ ex: jobs in a batch system; processes from different users
 - ♦ *cooperative* – if they share information or explicitly communicate
 - ♦ the *sharing* requires a *common address space*
 - ♦ *communication* can be done through a common address space or a *communication channel* connecting them

Concepts



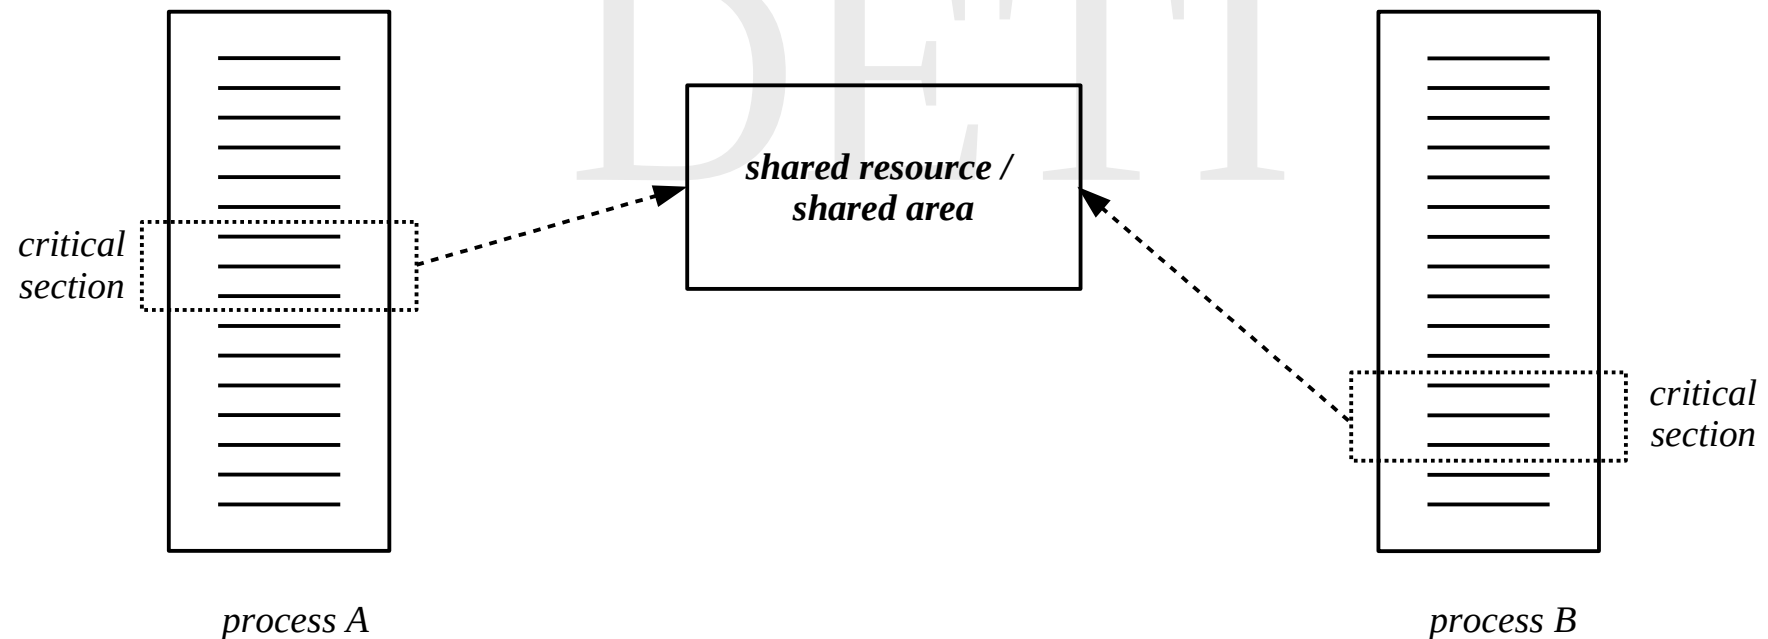
- *independent processs* competing for a resource
- is responsibility of the OS to guarantee that the assignment of resources to processes is done in a controlled way, such that no information lost occurs
- in general, this imposes that only one process can use the resource at a time -- **mutual exclusive access**



- *cooperative processes* sharing information or communicating
- is responsibility of the processes to guarantee that access to the shared area is done in a controlled way, such that no information lost occurs
- in general, this imposes that only one process can use the resource at a time -- **mutual exclusive access**
- the communication channel is typically a system resource; so processes compete for it

Concepts

- Having access to a resource or to a shared area, actually means executing the code that do the access
- This code, because needs to avoid **race conditions** (that result in lost of information), is called **critical section**



Concepts

- ♦ Mutual exclusion in the access to a resource or shared area can result in:
 - ♦ **deadlock** – when two or more processes are waiting forever for access to their respective critical section, waiting for events that can be demonstrated will never happen
 - ♦ operations are blocked
 - ♦ **starvation** – when one or more processes compete for access to a critical section and, due to a conjunction of circumstances in which new processes that exceed them continually arise, access is successively deferred
 - ♦ operations are continuously postponed

Access to a resource

/ processes competing for a resource - $p = 0, 1, \dots, N-1$ */*

void main (**unsigned int** p)

{

forever

 {

 do_something();

 access_resource(p);

 do_something_else();

 }

}

enter_critical_section(p);

use_resource();

leave_critical_section(p);

← **critical
section**

Access to a shared area

```
/* shared data structure */
shared DATA d;
/* processes sharing data -  $p = 0, 1, \dots, N-1$  */
void main (unsigned int p)
{
    forever
    {
        do_something();
        access_shared_area(p);
        do_something_else();
    }
}
```

Diagram illustrating the critical section access:

- The code block `access_shared_area(p);` is enclosed in a box, indicating it is the critical section.
- The box is flanked by `enter_critical_section(p);` and `leave_critical_section(p);` in blue text.
- An arrow points from the text **critical section** to the box.

Producer / consumer relationship

/ communicating data structure: FIFO of fixed size */*

shared FIFO fifo;

/ producer processes - $p = 0, 1, \dots, N-1$ */*

void main (**unsigned int** p)

{

DATA val;
bool done;

forever

{

produce_data(&val);
done = **false**;

do

{

enter_critical_section(p);

if (fifo.notFull())

{

fifo.insert(val);
done = **true**;

}

leave_critical_section(p);

} **while** (!done);
do_something_else();

}

}

DETI

← ***critical section***

Producer / consumer relationship

/ communicating data structure: FIFO of fixed size */*

shared FIFO fifo;

/ consumer processes - $p = 0, 1, \dots, M-1$ */*

void main (**unsigned int** p)

{

DATA val;
bool done;

forever

{

done = **false**;

do

{

enter_critical_section(p);

if (fifo.notEmpty())

{

fifo.retrieve(&val);

done = **true**;

}

leave_critical_section(p);

} **while** (!done);

consume_data(val);

do_something_else();

}

}

← ***critical section***

Access to a critical section

- ♦ Requirements that should be observed in accessing a critical section:
 - ♦ effective mutual exclusion – access to the critical sections associated with the same resource, or shared area, can only be allowed to one process at a time, among all processes that compete for access
 - ♦ independence on the number of intervening processes or on their relative speed of execution
 - ♦ a process outside the critical section can not prevent another from entering there
 - ♦ a process requiring access to the critical section should not have to wait indefinitely
 - ♦ length of stay inside a critical section should be necessarily finite

Type of solutions

- ♦ In general, a memory location is used to control access to the critical region
- ♦ **software solutions** – solutions that are based on the typical instructions used to access memory location
 - ♦ read and write are done by different instructions
- ♦ **hardware solutions** – solutions that are based on special instructions to access the memory location
 - ♦ these instructions allow to read and then write a memory location in an atomic way

Strict alternation

```
/* control data structure */
#define R      ...      /* process id = 0, 1, ..., R-1 */

shared unsigned int access_turn = 0;

void enter_critical_section(unsigned int own_pid)
{
    while (own_pid != access_turn);
}

void leave_critical_section(unsigned int own_pid)
{
    if (own_pid == access_turn)
        access_turn = (access_turn + 1) % R;
}
```

Strict alternation

- ♦ Not a valid solution
 - ♦ Dependence on the relative speed of execution of the intervening processes
 - ♦ The process with less accesses imposes its rhythm to the others
 - ♦ A process outside the critical section can prevent another from entering there
 - ♦ If it is not its turn, a process has to wait, even if no one else wants to enter

Constructing a solution

```
/* control data structure */
#define R      2          /* process id = 0, 1 */

shared bool is_in[R] = {false, false};

void enter_critical_section(unsigned int own_pid)
{
    unsigned int other_pid = 1 - own_pid;

    while (is_in[other_pid]);
    is_in[own_pid] = true;
}

void leave_critical_section(unsigned int own_pid)
{
    is_in[own_pid] = false;
}
```

Constructing a solution

- ♦ Not a valid solution
 - ♦ Mutual exclusion is not guaranteed
 - ♦ Assume that:
 - ♦ P_0 enters `enter_critical_section` and tests `is_in[1]`, which is *false*
 - ♦ P_1 enters `enter_critical_section` and tests `is_in[0]`, which is *false*
 - ♦ P_1 changes `is_in[0]` to *true* and enters its critical section
 - ♦ P_0 changes `is_in[1]` to *true* and enters its critical section
 - ♦ Thus, both processes enter the critical sections
- ♦ It seems that the failure is a result of testing first the other's control variable and then change its own variable

Constructing a solution

```
/* control data structure */
#define R 2 /* process pid = 0, 1 */

shared bool want_enter[R] = {false, false};

void enter_critical_section (unsigned int own_pid)
{
    unsigned int other_pid = 1 - own_pid;

    want_enter[own_pid] = true;
    while (want_enter[other_pid]);
}

void leave_critical_section (unsigned int own_pid)
{
    want_enter[own_pid] = false;
}
```


Constructing a solution

- ♦ Not a valid solution
 - ♦ Mutual exclusion is guaranteed, but deadlock can occur
 - ♦ Assume that:
 - ♦ P_0 enters `enter_critical_section` and sets `want_enter[0]` to *true*
 - ♦ P_1 enters `enter_critical_section` and sets `want_enter[1]` to *true*
 - ♦ P_1 tests `want_enter[0]` and, because it is *true*, keeps waiting to enter its critical section
 - ♦ P_0 tests `want_enter[1]` and, because it is *true*, keeps waiting to enter its critical section
 - ♦ Thus, both processes enter in deadlock
- ♦ To solve the deadlock at least one of the processes have to go back

Constructing a solution

```
/* control data structure */
#define R 2 /* process id = 0, 1 */

shared bool want_enter[R] = {false, false};

void enter_critical_section(unsigned int own_pid)
{
    unsigned int other_pid = 1 - own_pid;

    want_enter[own_pid] = true;
    while (want_enter[other_pid])
    {
        want_enter[own_pid] = false;
        random_delay();
        want_enter[own_pid] = true;
    }
}

void leave_critical_section(unsigned int own_pid)
{
    want_enter[own_pid] = false;
}
```

Constructing a solution

- ♦ An almost valid solution
 - ♦ The Ethernet protocol uses a similar approach to control access to the communication medium
- ♦ But, still not completely valid
 - ♦ Even if unlikely, deadlock and starvation can still be present
- ♦ The solution needs to be deterministic, not random

Dekker algorithm (1965)

```
#define R    2    /* process id = 0, 1 */
```

```
shared bool want_enter[R] = {false, false};
```

```
shared uint p_w_priority = 0;
```

```
void enter_critical_section(uint own_pid)
```

```
{
```

```
    uint other_pid = 1 - own_pid;
```

```
    want_enter[own_pid] = true;
```

```
    while (want_enter[other_pid])
```

```
    {
```

```
        if (own_pid != p_w_priority)
```

```
        {
```

```
            want_enter[own_pid] = false;
```

```
            while (own_pid != p_w_priority);
```

```
            want_enter[own_pid] = true;
```

```
        }
```

```
    }
```

```
}
```

```
void leave_critical_section(uint own_pid)
```

```
{
```

```
    uint other_pid = 1 - own_pid;
```

```
    p_w_priority = other_pid;
```

```
    want_enter[own_pid] = false;
```

```
}
```

Dekker algorithm (1965)

- ♦ The algorithm uses an alternation mechanism to solve the conflict
- ♦ Mutual exclusion in the access to the critical section is guaranteed
- ♦ Deadlock and starvation are not present
- ♦ No assumptions are done in the relative speed of the intervening processes
- ♦ However, it can not be generalized to more than 2 processes, satisfying all the requirements

Dijkstra algorithm (1966)

```
#define R    ...    /* process id = 0, 1, ..., R-1 */
```

```
shared uint want_enter[R] = {NO, NO, ... , NO};  
shared uint p_w_priority = 0;
```

```
void enter_critical_section(uint own_pid)  
{  
    uint n;  
    do  
    {  
        want_enter[own_pid] = WANT;  
        while (own_pid != p_w_priority)  
            if (want_enter[p_w_priority] == NO)  
                p_w_priority = own_pid;  
        want_enter[own_pid] = DECIDED;  
        for (n = 0; n < R; n++)  
            if (n != own_pid && want_enter[n] == DECIDED)  
                break;  
    } while (n < R);  
}
```

```
void leave_critical_section(uint own_pid)  
{  
    p_w_priority = (own_pid + 1) % R;  
    want_enter[own_pid] = NO;  
}
```

- Can suffer from starvation

Peterson algorithm (1981)

```
#define R    2    /* process id = 0, 1 */

shared bool want_enter[R] = {false, false};
shared uint last;      quem foi o ultimo processo a correr este codigo

void enter_critical_section(uint own_pid)
{
    uint other_pid = 1 - own_pid;
    want_enter[own_pid] = true;
    last = own_pid;
    while ((want_enter[other_pid]) && (last == own_pid));
}

void leave_critical_section(uint own_pid)
{
    want_enter[own_pid] = false;
}
```

Peterson algorithm (1981)

- ♦ The Peterson algorithm uses the order of arrival to solve conflicts
 - ♦ Each process has to write its ID in a shared variable (last)
 - ♦ The subsequent reading allows to determine which was the last one
- ♦ It is a valid solution
 - ♦ Guarantees mutual exclusion
 - ♦ Avoids deadlock and starvation
 - ♦ Make no assumption about the relative speed of intervening processes
- ♦ Can be generalized to more than processes
 - ♦ The general solution is similar to a waiting queue

Generalized Peterson algorithm (1981)

```
#define R    ...    /* process id = 0, 1, ..., R-1 */

shared int want_enter[R] = {-1, -1, ... , -1};
shared int last[R-1];

void enter_critical_section(uint own_pid)
{
    for (uint i = 0; i < R-1; i++)
    {
        want_enter[own_pid] = i;
        last[i] = own_pid;
        do
        {
            test = false;
            for (uint j = 0; j < R; j++)
                if (j != own_pid)
                    test = test || (want_enter[j] >= i);
        } while (test && (last[i] == own_pid));
    }
}

void leave_critical_section(int own_pid)
{
    want_enter[own_pid] = -1;
}
```

Hardware solutions - disabling interrupts

Uniprocessor computational system

- ♦ The switching of processes, in a multiprogrammed environment, is always caused by an external device:
 - ♦ *real time clock (RTC)* – causing the time-out transition in preemptive systems
 - ♦ *device controller* – can cause the preemp transitions in case of wake up of a higher priority process
 - ♦ In any case, interruptions of the processor
- ♦ Thus, access in mutual exclusion can be implemented disabling interrupts
- ♦ Only valid in kernel
 - ♦ Malicious or buggy code can completely block the system

Multiprocessor computational system

- ♦ Disabling interrupts in one processor has no effect

Hardware solutions - special instructions

```
shared bool flag = false;
```

```
bool test_and_set(bool * flag)
{
    bool prev = *flag;
    *flag = true;
    return prev;
}
```

```
void lock(bool * flag)
{
    while (test_and_set(flag));
}
```

```
void unlock(bool * flag)
{
    *flag = false;
}
```

- The **test_and_set** function, if implemented atomically (without interruptions), can be used to construct the lock (enter critical section) primitive
- In the instruction set of some of the current processors, there is an atomic instruction implementing that behavior
- Surprisingly, it is often called TAS (test and set)

Hardware solutions - special instructions

```
shared int value = 0;
```

```
int compare_and_swap(int * value, int expected, int new_value)
{
    int v = *value;
    if (*value == expected)
        *value = new_value;
    return v;
}
```

- The **compare_and_swap** function, if implemented atomically (without interruptions), can be used to construct the lock (enter critical section) primitive

```
void lock(int * flag)
{
    while (compare_and_swap(&flag, 0, 1) != 0);
}
```

```
void unlock(bool * flag)
{
    *flag = 0;
}
```

- In the instruction set of some of the current processors, there is an atomic instruction implementing that behavior
- In some instruction sets, there is a **compare_and_set** variant that returns a bool

Busy waiting

- ♦ The previous solutions suffer from **busy waiting** – the lock primitive is in the active state (using the CPU) while waiting
 - ♦ They are often referred as **spinlocks**, as the process spins around the variable while waiting for access
- ♦ In uniprocessor systems, busy waiting is unwanted, as there is
 - ♦ **loss of efficiency** – the time quantum of a process can be used for nothing
 - ♦ **risk of deadlock** – if a higher priority process calls lock while a lower priority process is inside its critical section, none of them can proceed
- ♦ In multiprocessor system with shared memory, busy waiting can be less critical
 - ♦ switching processes cost time, that can be higher than the time spent by the other process inside its critical section

block and wake_up

- ♦ In general, at least in uniprocessor systems, there is the requirement of blocking a process while it is waiting for entering its critical section

```
#define R    ...    /* process id = 0, 1, ..., R-1 */
```

```
shared unsigned int access = 1;
```

```
void enter_critical_section(unsigned int own_pid)
```

```
{
```

```
    if (access == 0) block(own_pid);  
    else access -= 1;
```

```
}
```

→ { ***atomic operation***
(can not be interrupted)

```
void leave_critical_section(unsigned int own_pid)
```

```
{
```

```
    if (there_are_blocked_processes) wake_up_one();  
    else access += 1;
```

```
}
```

→ { ***atomic operation***
(can not be interrupted)

- ♦ Atomic operations are still required
- ♦ Note that `access` is an integer, not a boolean

Semaphores

- A *semaphore* is a synchronization mechanism, defined by a data type plus two atomic operations, *down* and *up*
 - The operations are also referred to as *wait* and *signal/post*, respectively

- Data type:

```
typedef struct
{
    unsigned int val;      /* can not be negative */
    PROCESS *queue;       /* queue of waiting blocked processes */
} SEMAPHORE;
```

- Operations:
 - *down* – block process if *val* is zero; decrement *val* otherwise
 - *up* – if *queue* is not empty, wake up one process (accordingly to a given policy); increment *val* otherwise
- Note that *val* can only be manipulated through these operations

Semaphores

```
/* array of semaphores defined in kernel */
#define R    ... /* semid = 0, 1, ..., R-1 */

static SEMAPHORE sem[R];

void sem_down(unsigned int semid)
{
    disable_interruptions;
    if (sem[semid].val == 0)
        block_on_sem(getpid(), semid);
    else
        sem[semid].val -= 1;
    enable_interruptions;
}

void sem_up(unsigned int semid)
{
    disable_interruptions;
    if (sem[semid].queue != NULL)
        wake_up_one_on_sem(semid);
    else
        sem[semid].val += 1;
    enable_interruptions;
}
```

- This implementation is typical of uniprocessor systems. *Why?*

- Semaphores can be binary or not binary
- How to implement *mutual exclusion* using semaphores?

Bounded-buffer problem

```
shared FIFO fifo;    /* fixed-size FIFO memory */
```

```
/* producers -  $p = 0, 1, \dots, N-1$  */
```

```
void producer(unsigned int p)
```

```
{
```

```
    DATA data;
```

```
    forever
```

```
    {
```

```
        produce_data(&data);
```

```
        bool done = false;
```

```
        do
```

```
        {
```

```
            lock(p);
```

```
            if (fifo.notFull())
```

```
            {
```

```
                fifo.insert(data);
```

```
                done = true;
```

```
            }
```

```
            unlock(p);
```

```
        } while (!done);
```

```
        do_something_else();
```

```
    }
```

```
}
```

```
/* consumers -  $c = 0, 1, \dots, M-1$  */
```

```
void consumer(unsigned int c)
```

```
{
```

```
    DATA data;
```

```
    forever
```

```
    {
```

```
        bool done = false;
```

```
        do
```

```
        {
```

```
            lock(c);
```

```
            if (fifo.notEmpty())
```

```
            {
```

```
                fifo.retrieve(&data);
```

```
                done = true;
```

```
            }
```

```
            unlock(c);
```

```
        } while (!done);
```

```
        consume_data(data);
```

```
        do_something_else();
```

```
    }
```

```
}
```

- ♦ How to implement using semaphores?
 - ♦ Guaranteeing mutual exclusion and absence of busy waiting

Solving the bounded-buffer problem using semaphores

```
shared FIFO fifo;      /* fixed-size FIFO memory */
shared sem access;     /* semaphore to control mutual exclusion */
shared sem nslots;     /* semaphore to control number of available slots*/
shared sem nitems;     /* semaphore to control number of available items */
```

```
/* producers -  $p = 0, 1, \dots, N-1$  */
```

```
void producer(unsigned int p)
{
    DATA val;

    forever
    {
        produce_data(&val);
        sem_down(nslots);
        sem_down(access);
        fifo.insert(val);
        sem_up(access);
        sem_up(nitems);
        do_something_else();
    }
}
```

```
/* consumers -  $c = 0, 1, \dots, M-1$  */
```

```
void consumer(unsigned int c)
{
    DATA val;

    forever
    {
        sem_down(nitems);
        sem_down(access);
        fifo.retrieve(&val);
        sem_up(access);
        sem_up(nslots);
        consume_data(val);
        do_something_else();
    }
}
```

- `fifo.empty()` and `fifo.full()` are not necessary. *Why?*
- What are the initial values of the semaphores?

Wrong solution of the bounded-buffer problem

```
shared FIFO fifo;      /* fixed-size FIFO memory */
shared sem access;     /* semaphore to control mutual exclusion */
shared sem nslots;     /* semaphore to control number of available slots*/
shared sem nitems;     /* semaphore to control number of available items */
```

```
/* producers -  $p = 0, 1, \dots, N-1$  */
```

```
void producer(unsigned int p)
{
    DATA data;
    forever
    {
        produce_data(&data);
        sem_down(access);
        sem_down(nslots);
        fifo.insert(data);
        sem_up(access);
        sem_up(nitems);
        do_something_else();
    }
}
```

```
/* consumers -  $c = 0, 1, \dots, M-1$  */
```

```
void consumer(unsigned int c)
{
    DATA data;
    forever
    {
        sem_down(nitems);
        sem_down(access);
        fifo.retrieve(&data);
        sem_up(access);
        sem_up(nslots);
        consume_data(data);
        do_something_else();
    }
}
```

- What is wrong with this solution?

Analysis of semaphores

- ♦ Concorrent solutions based on semaphores have advantages and disadvantages
- ♦ *Advantages:*
 - ♦ *support at the operating system level* – operations on semaphores are implemented by the kernel and made available to programmers as *system calls*
 - ♦ *general* – they are low level contructions and so they are versatile, being able to be used in any type of solution
- ♦ *Disadvantages*
 - ♦ *specialized knowledge* – the programmer must be aware of concorrent programming principles, as race conditions or deadlock can be easily introduced
 - ♦ See the previous example, as an illustration of this

Monitors

- ♦ A problem with semaphores is that they are used both to implement mutual exclusion and to synchronize processes
- ♦ Being low level primitives, they are applied in a bottom-up perspective
 - ♦ if required conditions are not satisfied, processes are blocked before they enter their critical sections
 - ♦ this approach is prone to errors, mainly in complex situations, as synchronization points can be scattered throughout the program
- ♦ A higher level approach should followed a top-down perspective
 - ♦ processes must first enter their critical regions and then block if pursuance conditions are not satisfied
- ♦ A solution is to introduce a (concurrent) construction at the programming language level that separately deals with mutual exclusion and synchronization
- ♦ A *monitor* is a synchronization mechanism, independently proposed by Hoare and Brinch Hansen, supported by a (concurrent) programming language
 - ♦ It is composed of an internal data structure, inicialization code and a number of accessing primitives

Monitors

monitor example

```
{
    /* internal shared data structure */
    DATA data;

    condition c; /* condition variable */

    /* access methods */
    method_1 (...)
    {
        ...
    }

    method_2 (...)
    {
        ...
    }


    ...

    /* initialization code */
    ...
}
```

- An application is seen as a set of threads that compete to access the shared data structure
- This shared data can only be accessed through the access methods
- Every method is executed in mutual exclusion
- If a thread calls an access method while another thread is inside another access method, its execution is blocked until the other leaves
- Synchronization between threads is possible through *condition variables*
- Two operation on them are possible:
 - *wait* – the thread is blocked and put outside the monitor
 - *signal* – if there are threads blocked, one is waked up. *Which one?*

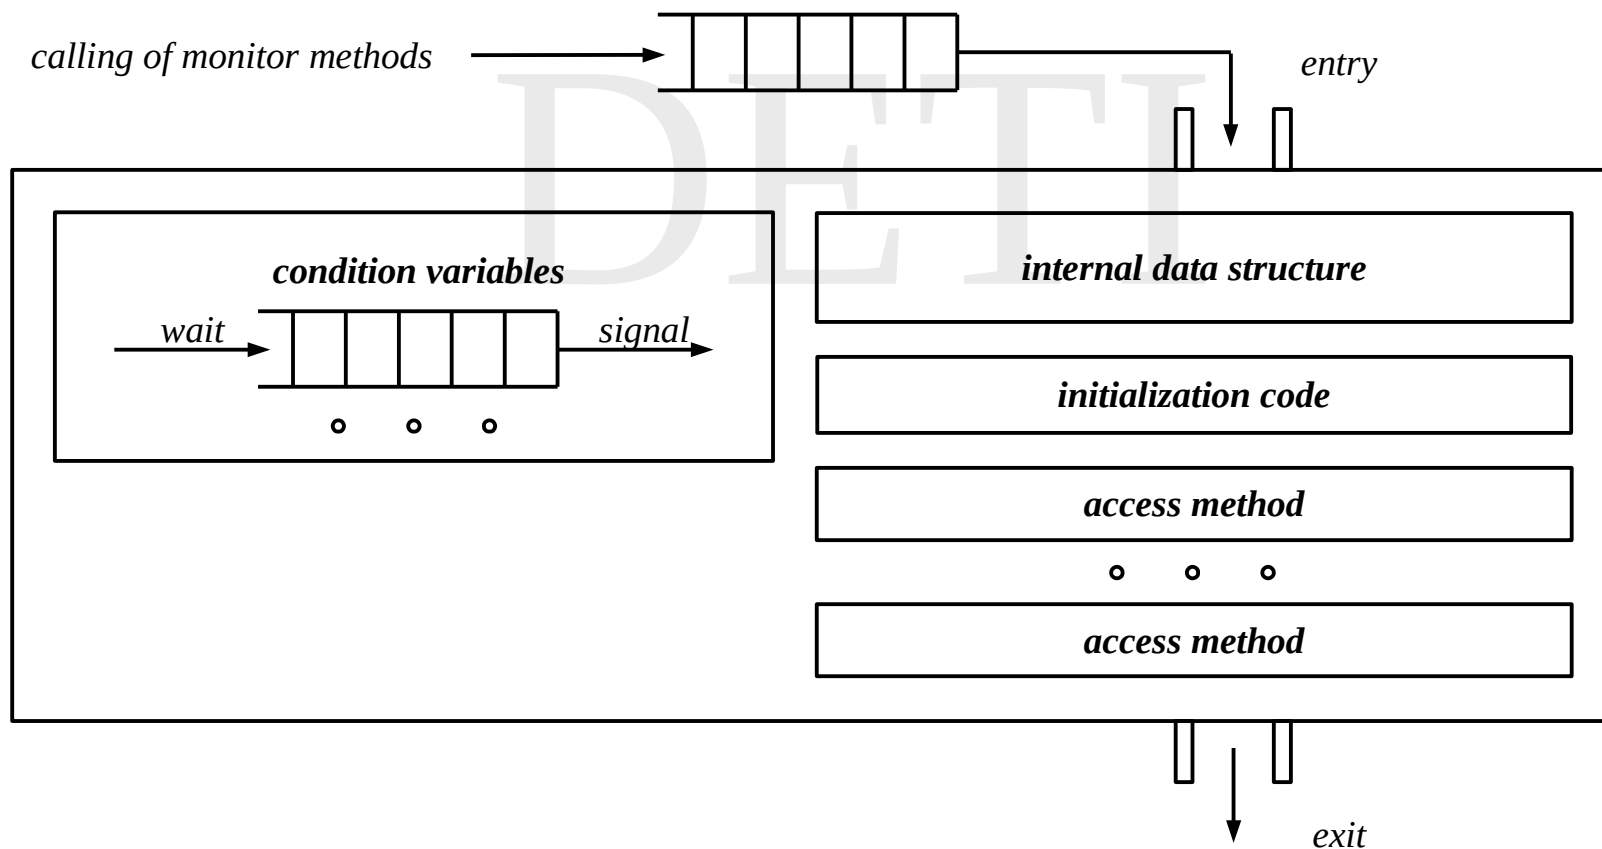
Hoare monitor

- What to do when *signal* occurs?
- *Hoare monitor* – the thread calling *signal* is put out of the monitor, so the just waked up thread can proceed
 - quite general, but its implementation requires a stack where the blocked thread is put



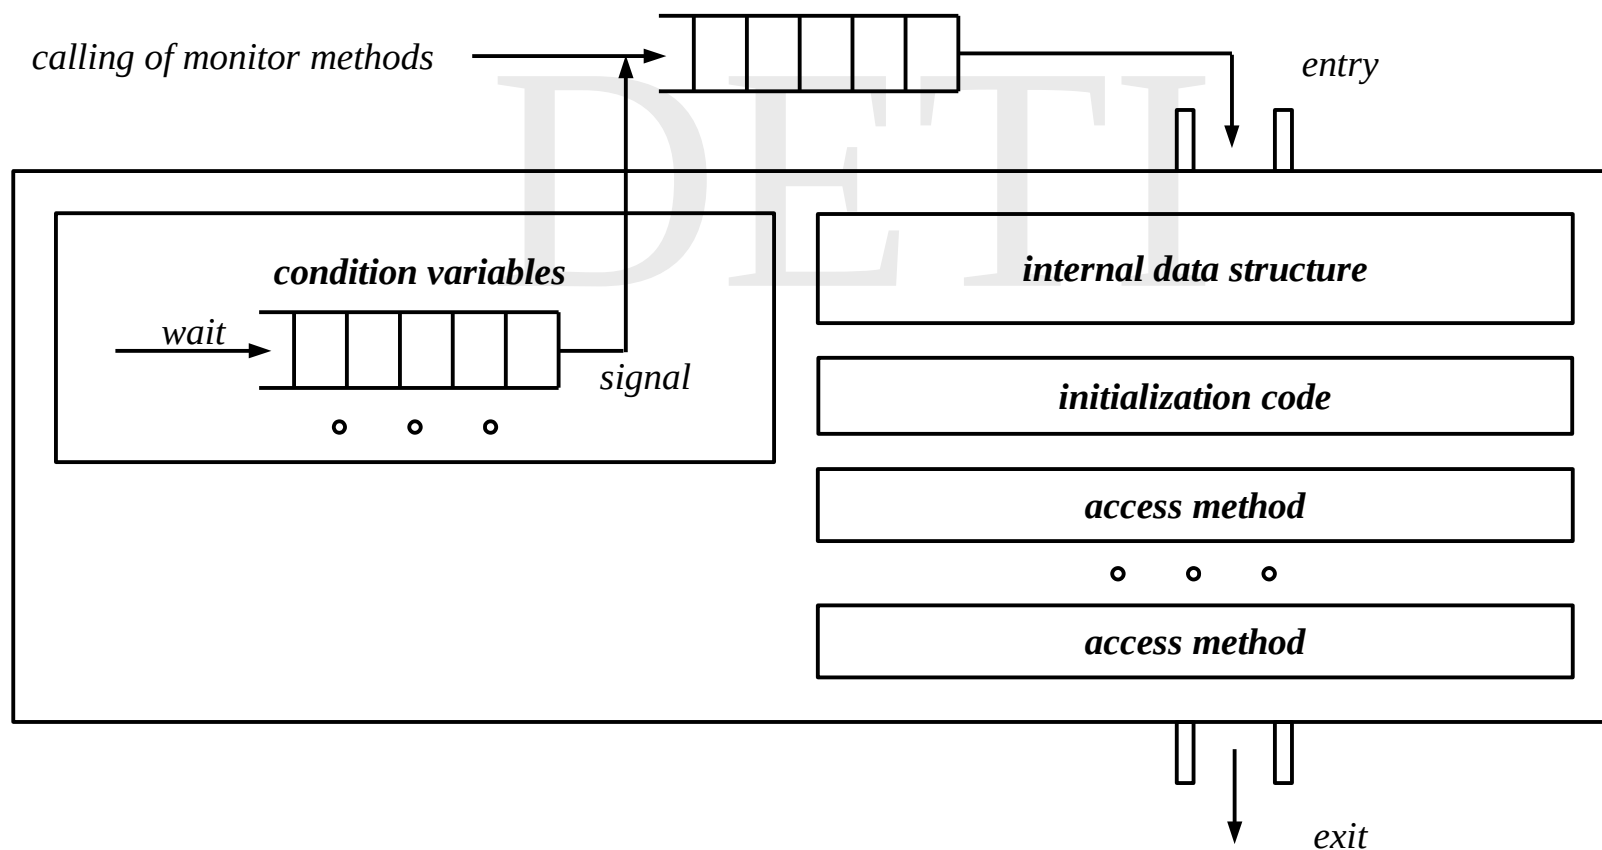
Brinch Hansen monitor

- ♦ *Brinch Hansen monitor* – the thread calling *signal* immediately leaves the monitor (*signal* is the last instruction of the monitor method)
- ♦ easy to implement, but quite restrictive (only one signal allowed in a method)



Lampson / Redell monitors

- ♦ *Lampson / Redell monitor* – the thread calling signal continues its execution and the just waked up thread is kept outside the monitor, competing for access
- ♦ easy to implement, but can cause starvation



Solving the bounded-buffer problem using monitors

```
shared FIFO fifo;      /* fixed-size FIFO memory */
shared mutex access;   /* mutex to control mutual exclusion */
shared cond nslots;    /* condition variable to control availability of slots*/
shared cond nitems;    /* condition variable to control availability of items */
```

```
/* producers -  $p = 0, 1, \dots, N-1$  */
```

```
void producer(unsigned int p)
```

```
{
    DATA data;
    forever
    {
        produce_data(&data);
        lock(access);
        if/while (fifo.isFull())
        {
            wait(nslots, access);
        }
        fifo.insert(data);
        unlock(access);
        signal(nitems);
        do_something_else();
    }
}
```

```
/* consumers -  $c = 0, 1, \dots, M-1$  */
```

```
void consumer(unsigned int c)
```

```
{
    DATA data;
    forever
    {
        lock(access);
        if/while (fifo.isEmpty())
        {
            wait(nitems, access);
        }
        fifo.retrieve(&data);
        unlock(access);
        signal(nslots);
        consume_data(data);
        do_something_else();
    }
}
```

- `fifo.empty()` and `fifo.full()` are now necessary. Why?
- What is the initial value of the mutex?

Message-passing

- ♦ Processes can communicate exchanging messages
 - ♦ A general communication mechanism, not requiring shared memory
 - ♦ Valid for uniprocessor and multiprocessor systems
- ♦ Two operation are required:
 - ♦ **send** and **receive**
- ♦ A communication link is required:
 - ♦ There are different logical ways of implementing it
 - ♦ Direct or indirect (through mailboxes or ports) addressing
 - ♦ Synchronous or asynchronous communication
 - ♦ Automatic or explicit buffering

Message-passing - direct vs. indirect

- ♦ Symmetric direct communication
 - ♦ A process that wants to communicate must explicitly name the recipient or sender/receiver
 - ♦ **send**(P, message) - send message to process P
 - ♦ **receive**(P, message) - receive message from process P
 - ♦ A communication link in this scheme has the following properties:
 - ♦ it is established automatically between a pair of communicating processes
 - ♦ it is associated with exactly two processes
 - ♦ between a pair of processes there exist exactly one link
- ♦ Asymmetric direct communication
 - ♦ Only the sender must explicitly name the recipient
 - ♦ **send**(P, message) - send message to process P
 - ♦ **receive**(id, message) - receive message from any process

Message-passing - direct vs. indirect

- ♦ Indirect communication
 - ♦ The messages are sent and received from mailboxes, or ports
 - ♦ **send**(M, message) - send message to mailbox M
 - ♦ **receive**(M, message) - receive message from mailbox M
 - ♦ A communication link in this scheme has the following properties:
 - ♦ it is only established if the pair of communicating processes has a shared mailbox
 - ♦ it may be associated with more than two processes
 - ♦ between a pair of processes there may exist more than one link (a mailbox per each)
 - ♦ The problem of two or more processes trying to receive a message from the same mailbox:
 - ♦ Is it allowed?
 - ♦ If allowed, which one will succeed?

Message-passing - synchronization

- ♦ There are different design options for implementing `send` and `receive`
 - ♦ **Blocking send** - the sending process blocks until the message is received by the receiving process or by the mailbox
 - ♦ **Nonblocking send** - the sending process sends the message and resumes operation.
 - ♦ **Blocking receive** - the receiver blocks until a message is available
 - ♦ **Nonblocking receive** - the receiver retrieves either a valid message or the indication that no one exists
- ♦ Different combinations of `send` and `receive` are possible

Message-passing - buffering

- ♦ There are different design options for implementing the link supporting the communication
 - ♦ **Zero capacity** - there is no queue
 - ♦ the sender must block until the recipient receives the message
 - ♦ **Bounded capacity** - the queue has finite length
 - ♦ if the queue is full, the sender must block until space is available
 - ♦ **Unbounded capacity** - the queue has (potentially) infinite length

Solving the bounded-buffer problem using messages

```
/* producers -  $p = 0, 1, \dots, N-1$  */  
void producer(unsigned int p)  
{  
    DATA data;  
    MESSAGE msg;  
  
    forever  
    {  
        produce_data(&val);  
        make_message(msg, data);  
        send(msg);  
        do_something_else();  
    }  
}
```

```
/* consumers -  $c = 0, 1, \dots, M-1$  */  
void consumer(unsigned int c)  
{  
    DATA data;  
    MESSAGE msg;  
  
    forever  
    {  
        receive(msg);  
        extract_data(data, msg);  
        consume_data(data);  
        do_something_else();  
    }  
}
```


POSIX support for monitors

- ♦ Standard *POSIX*, *IEEE 1003.1c*, defines a programming interface (API) for the creation and synchronization of *threads*.
 - ♦ In unix, this interface is implemented by the *pthread* library
- ♦ It allows for the implementation of monitors in C/C++
 - ♦ Using mutexes and condition variables
 - ♦ Note that they are of the Lampson / Redell type
- ♦ Some of the available functions:
 - ♦ *pthread_create* – creates a new thread; similar to `fork`
 - ♦ *pthread_exit* – equivalent to `exit`
 - ♦ *pthread_join* – equivalent a `waitpid`
 - ♦ *pthread_self* – equivalent a `getpid()`
 - ♦ *pthread_mutex_** – manipulation of mutexes
 - ♦ *pthread_cond_** – manipulation of condition variables
 - ♦ *pthread_once* – initialization

Semaphores in Unix/Linux

- ♦ System V semaphores
 - ♦ creation: `semget`
 - ♦ down and up: `semop`
 - ♦ other operations: `semctl`
- ♦ POSIX semaphores
 - ♦ down and up
 - ♦ `sem_wait`, `sem_trywait`, `sem_timedwait`, `sem_post`
 - ♦ Two types: named and unnamed semaphores
 - ♦ Named semaphores
 - ♦ `sem_open`, `sem_close`, `sem_unlink`
 - ♦ created in a virtual filesystem (e.g., `/dev/sem`)
 - ♦ unnamed semaphores - memory based
 - ♦ `sem_init`, `sem_destroy`
 - ♦ execute `man sem_overview` for an overview

Message-passing in Unix/Linux

- ♦ System V implementation
 - ♦ Defines a message queue where messages of different types (a positive integer) can be stored
 - ♦ The send operation blocks if space is not available
 - ♦ The receive operation has an argument to specify the type of message to receive: a given type, any type or a range of type
 - ♦ The oldest message of given type(s) is retrieved
 - ♦ Can be blocking or nonblocking
 - ♦ see system calls: `msgget`, `msgsnd`, `msgrcv`, and `msgctl`
- ♦ POSIX message queue
 - ♦ Defines a priority queue
 - ♦ The send operation blocks if space is not available
 - ♦ The receive operation removes the oldest message with the highest priority
 - ♦ Can be blocking or nonblocking
 - ♦ see functions: `mq_open`, `mq_send`, `mq_receive`, ...

Shared memory in Unix/Linux

- ♦ Address spaces are independent
- ♦ But address spaces are virtual
- ♦ The same physical region can be mapped into two or more virtual regions
- ♦ This is managed as a resource by the operating system
- ♦ System V shared memory
 - ♦ creation - `shmget`
 - ♦ mapping - `shmat`, `shmdt`
 - ♦ other operations – `shmctl`
- ♦ POSIX shared memory
 - ♦ creation - `shm_open`, `ftruncate`
 - ♦ mapping - `mmap`, `munmap`
 - ♦ other operations - `close`, `shm_unlink`, `fchmod`, ...

Bibliography

Operating Systems Concepts, Silberschatz, Galvin, Gagne, John Wiley & Sons, 9th Ed

- Chapter 3: *Processes* (section 3.4)
- Chapter 5: *Process synchronization* (sections 5.1 to 5.8)

Modern Operating Systems; Tanenbaum & Bos; Prentice-Hall International Editions, 4rd Ed

- Chapter 2: *Processes and Threads* (sections 2.3 and 2.5)

Operating Systems, Stallings, Prentice-Hall International Editions, 7th Ed

- Chapter 5: *Concurrency: mutual exclusion and synchronization* (sections 5.1 to 5.5)